

A SHARED PLACE FOR
CARE

Less to carry. More care to give.

 Carestead GEN AI ACADEMY · TEAM PROJECT

Running local prototype · synthetic
care records

A caregiver
cognitive-load
agent that turns a

Care happens across people. The mental checklist stays with one.



A FAMILIAR COORDINATION GAP

“I thought you were taking Alex tomorrow.”

The appointment changes

The ride still points to the old plan.

Ownership is unclear

Two people assume the other has it.

Maya reconstructs

Small gaps can disrupt essential care

A missed pickup or unconfirmed ride can leave an important responsibility unfinished.

Every handover costs attention

The next caregiver needs the current situation, open issues, and reliable contacts.

More reminders are only part of the answer

Problem framing: [team project document](#) · illustrative scenario

One person. One current care picture.

Maya can hand over context, priorities, and responsibilities together.

A foundation for each recipient

Profile, communication and mobility needs, support contacts, and a versioned care plan.

A concise, live handover

Latest event, open risks, pending approvals, next tasks, and facts to verify. Copy the brief for a handover.

A plan that



Carestead

KNOW THE
PERSON ·
UNDERSTAND
TODAY

Implementation: 03
app/page.tsx · /
app/api/state/route.ts · 07
db/schema.ts

From “What needs attention?” to a reviewed action.



Carestead

ASK · REVIEW ·
APPROVE

Ask in the moment

Recipient-specific text chat returns a summary with expandable supporting evidence.

Prepare a concrete change

Reschedule an appointment, assign a ride, add a responsibility, run a check, or draft an in-app update.

Approve, then record the result



Implementation: `app/api/chat/route.ts` ·
`components/care-chat.tsx`

04 /
07



Structured
state
grounds
the agent.
Human
approval
controls
the write.

1 .
OBSERVE

**Caregiver
input**

Task
updates,
chat
requests, or
a care
check.

2 .
RETRIEVE

**Authorized
context**

Profile, tasks,
events, trusted
facts, risks,
and contacts.

3 .
DECIDE

**Rules +
typed
tools**

Return a
grounded
answer or
propose a
supported
action.

4 .
REVIEW

**Policy +
approval**

Recheck
role and
consent.
Execute
approved
writes only.

5 .
RECORD

**State +
trace**

Save the
change,
evidence,
tool
outcome,
and audit
entry.

**React /
Vinext →
API →
Cloudflare
D1**

One care-
state agent.
Structured
retrieval
supplies
grounded
context;
deterministic
rules and
intent
matching
drive
decisions.

No separate
vector
database,
Mem0
service, or
LLM in the
current

**Exp
de
bra**

Que
ans
evic
Miss
con
requ
clari
/ no
Prop
acti
app
canc
Den
peri
or c
→ b



Carestead

MEASURE
THE
DECISION
PATH

A measurable baseline. Visible failures to improve.

16 curated
synthetic
scenarios
across
seven
categories
•
benchmark
v1.0

76%

Evidence retrieval

Expected evidence strings matched

81%

Severity accuracy

Risk level matches ground truth

94%

Policy agreement

Approval flag matches the label

75%

Action accuracy

Action class matches ground truth

12 /
16
cases
pass

Severity, approval, and action all match.

This pass count does not require perfect retrieval. These are regression scores, not clinical validation or measured

Four failures define the next work

- Unavailable transportation owner - reassignment
- Appointment overlap resolve conflict
- Unanswered check-in escalate
- Memory review r → flag s context

Useful support, with a clear boundary around autonomy.

Permission before action

Server-enforced recipient roles, explicit approval for chat writes, and approval-aware in-app notifications.

Control over personal data

Consent withdrawal, retention choices, owner export and confirmed deletion; validation, rate limits, and privacy-safe error metadata.

Coordination only

No diagnosis, prescribing, or replacement for clinical judgment or emergency services.

Verified in the local walkthrough

Grounded chat; no task change before approval; task updated after approval. Automated WCAG checks on four surfaces.

Next: prove caregiver value

Close the four benchmark gaps. Add expert-reviewed cases. Measure verified task completion, caregiver effort, and unnecessary alerts.

Guardrails reduce risk; they do not establish clinical readiness. Real-world completion and emergency handling need further validation.

Know what matters. Share the responsibility. Approve the next step.